

## 3.1 Guarantees for policy gradient methods

What kinds of problems are policy gradient methods good at solving?

For example, consider a very simple MDP in which